



Team# 7

Shubham Vadhera | Sagar Dafle | Jagmohan Singh | Pooja Yelure

MongoDB Design Philosophy

- Combining the critical capabilities of relational databases with the innovations of NoSQL technologies



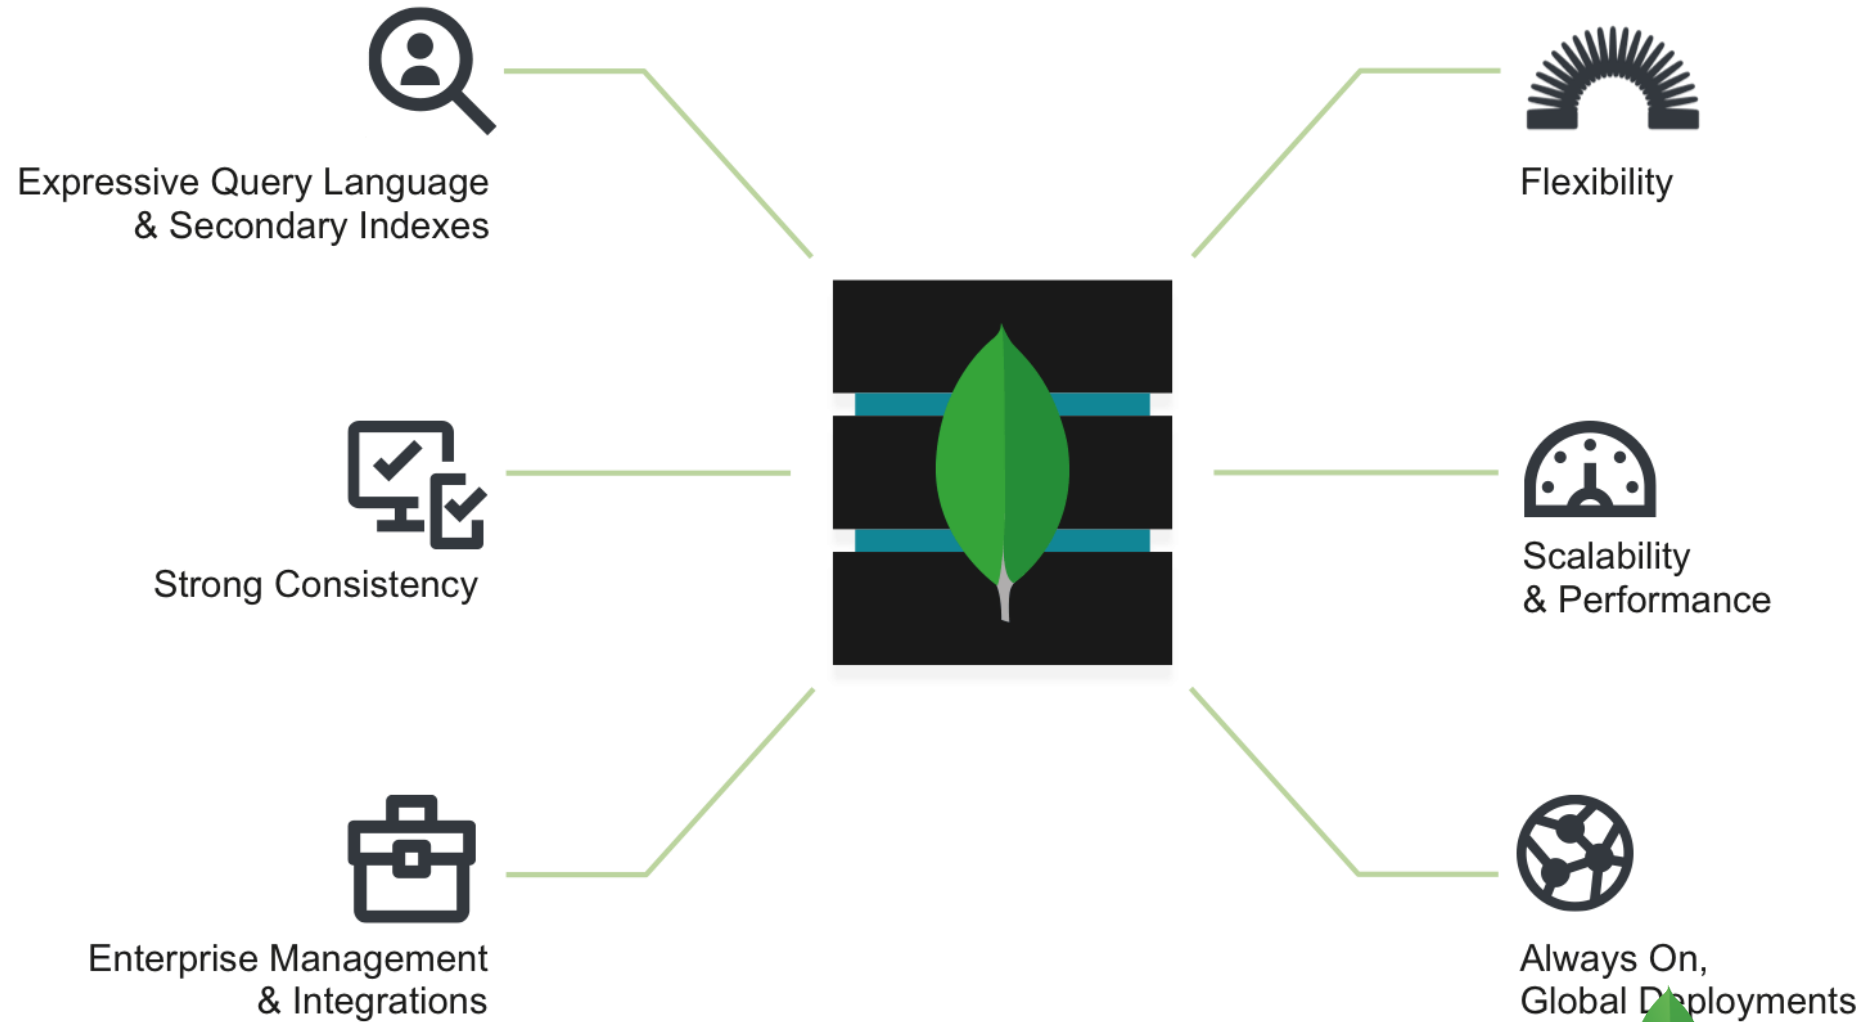
RELATIONAL



NoSQL



Mongo DB !



Traditional DB – Relational DB

- **Expressive query language** - Users should be able to access and manipulate their data in sophisticated ways
- **Strong consistency** - Applications should be able to immediately read what has been written to the database
- **Enterprise Management and Integrations** - A database should allow administrators to secure, monitor, automate, and integrate with their existing technology infrastructure

Modern DB – NoSQL DB

- **Flexible Data Model** - Easy to store and combine data of any structure
- **Scalability and Performance** - Enable almost unlimited growth with higher throughput and lower latency than relational databases.
- **Always-On Global Deployments** – Available systems across many nodes



MongoDB Data Model

- Data As Documents

id	user_name	email	age	city
1	Mark Hanks	mark@abc.com	25	Los Angeles
2	Richard Peter	richard@abc.com	31	Dallas



```
{
  "_id": ObjectId("5146bb52d8524270060001f3"),
  "age": 25,
  "city": "Los Angeles",
  "email": "mark@abc.com",
  "user_name": "Mark Hanks"
}
{
  "_id": ObjectId("5146bb52d8524270060001f2"),
  "age": 31,
  "city": "Dallas",
  "email": "richard@abc.com",
  "user_name": "Richard Peter"
}
```

- MongoDB stores data as documents in a binary representation called BSON (Binary JSON)
- BSON extends the popular JSON (JavaScript Object Notation) representation to include additional types such as int, long, date, and floating point
- Documents that tend to share a similar structure are organized as collections

MongoDB Data Model

- MongoDB documents are composed of field-and-value pairs and have the following structure:

```
{  
  field1: value1,  
  field2: value2,  
  field3: value3,  
  ...  
  fieldN: valueN  
}
```

```
var mydoc = {  
  _id: ObjectId("5099803df3f4948bd2f98391"),  
  name: { first: "Alan", last: "Turing" },  
  birth: new Date('Jun 23, 1912'),  
  death: new Date('Jun 07, 1954'),  
  contribs: [ "Turing machine", "Turing test", "Turingery" ],  
  views : NumberLong(1250000)  
}
```

MongoDB Data Model

- Dynamic Schema
 - Fields can vary from document to document
 - There is no need to declare the structure of documents to the system –documents are self describing
- Document Validation
 - Users can enforce checks on document structure, data types, data ranges and the presence of mandatory fields



Key Market Features (Pros) of MongoDB

- **High availability** (by replicating the data)
- **Scalability** (from a standalone server to distributed architectures of huge clusters). This allows us to shard our database transparently across all our shards. This increases the performance of our data processing.
- **Aggregation**: Batch data processing and aggregate calculations using native MongoDB operations.
- **Load Balancing**: Automatic data movement across *different shards* for load balancing. The balancer decides when to migrate the data and the destination Shard, so they are evenly distributed among all servers in the cluster.

<http://www.mongodbspain.com/en/2014/08/17/mongodb-characteristics-future/>



Why MONGO over other NoSQL DBs

- [MongoDB Management Service](#) (MMS) is a powerful **web tool** that allows us tracking our databases and our machines and also backing up our data.
- MMS tracks the database and hardware metrics for managing mongodb deployment.
- **Custom alerts:** Discover issues before your MongoDB instance will be affected.
- **Task Automation:** Simple launch and configuration of standalone MongoDB instances, replica sets or sharded clusters.

<http://www.mongodbspain.com/en/2014/08/17/mongodb-characteristics-future/>



CONS of Mongo DB

- **Less Flexibility** with querying (e.g. no JOINS)
- **Memory Usage**
MongoDB has the natural tendency to use up more memory because it has to store the key names within each document.
- **Concurrency Issues**
When you perform a write operation in MongoDB, it creates a lock on the entire database, not just the affected entries, and not just for a particular connection. This lock blocks not only other write operations, but also read operations.

<http://halls-of-valhalla.org/beta/articles/the-pros-and-cons-of-mongodb,45/>



Optimal use cases of Mongo

- Usually used when we need a **horizontally scalable performance** for high loads.
- Real-time analytics and high-speed logging, caching and high scalability.
- RDBMS replacement for **web applications**.
- With 1.5 million new classified ads posted every day, **Craigslist** must archive billions of records in many different formats, and must be able to query and report on these archives at runtime.

<https://www.mongodb.com/customers/craigslist>



CRUD Operations

- Read Operations: Queries are the core operations that return data in MongoDB.
- Cursors: Queries return iterable objects, called cursors, that hold the full result set.
- Write operations: Insert, update, or remove documents in MongoDB. Introduces data create and modify operations, their behavior, and performances.
- Atomicity and Transactions: Describes write operation atomicity in MongoDB.

CRUD Operations

- **Commands Used to Create Document:**

- `db.collection.insert()`
- `db.collection.insertOne();`
- `db.collection.insert();`

- **Example to Insert Query:** `db.bank_data.insertOne({"first_name": "John", "last_name": "Cena", "accounts": [{ "account_type": "Investment", "account_balance": 6123524974.110823463, "currency": "USD" }, { "account_type": "Savings", "account_balance": 132933272.569229168, "currency": "EURO" }]});`

- **Command Used to Read Documents:**

- `db.collection.find();`
- `Db.collection.findOne();`

- **Example for reading :** `db.bank_data.findOne({"last_name": "SMITH"});`

- **Projections:** In Mongo DB projection meaning is selecting only necessary data rather than selecting whole of the data of a document

- `$`
- `$elemMatch`
- `$limit`

- **Example:** `db.bank_data.findOne({"last_name": "Cena"}, {"first_name": 1, "_id": 0});`



Comparison Operators

<code>\$eq</code>	Matches values that are equal to a specified value.
<code>\$gt</code>	Matches values that are greater than a specified value.
<code>\$gte</code>	Matches values that are greater than or equal to a specified value.
<code>\$lt</code>	Matches values that are less than a specified value.
<code>\$lte</code>	Matches values that are less than or equal to a specified value.
<code>\$ne</code>	Matches all values that are not equal to a specified value.
<code>\$in</code>	Matches any of the values specified in an array.
<code>\$nin</code>	Matches none of the values specified in an array.

Example: `$eq:` `db.bank_data.find({"accounts.account_balance": {$eq:132933272.569229168}}).pretty()`
`$gte,$lte:` `db.bank_data.find({"accounts.account_balance": {$gte:8554996,$lte:9000000}}).pretty()[4];`

Logical Operators

Name	Description
<code>\$or</code>	Joins query clauses with a logical OR returns all documents that match the conditions of either clause.
<code>\$and</code>	Joins query clauses with a logical AND returns all documents that match the conditions of both clauses.
<code>\$not</code>	Inverts the effect of a query expression and returns documents that do <i>not</i> match the query expression.
<code>\$nor</code>	Joins query clauses with a logical NOR returns all documents that fail to match both clauses.

- Examples:
 - `$and db.bank_data.find({$and:[{"accounts.account_type":"Checking"}, {"accounts.account_balance":{$gte:8554996,$lte:9000000}}]}).pretty()[4];`
 - `$or db.bank_data.find({$and:[{"accounts.account_type":"Investment"}, {"accounts.account_type":"Savings"}]}).pretty()[4];`

Array and Regex Operator

Name	Description
<code>\$all</code>	Matches arrays that contain all elements specified in the query.
<code>\$elemMatch</code>	Selects documents if element in the array field matches all the specified <code>\$elemMatch</code> conditions.
<code>\$size</code>	Selects documents if the array field is a specified size.

Example `$elemMatch`: `db.bank_data.find({last_name: "SMITH", "accounts.account_type": "Savings" }, { first_name: 1, last_name: 1, accounts: { $elemMatch : { 'account_type' : 'Savings' } } }).pretty();`

\$regex: Provides regular expression capabilities for pattern matching *strings* in queries. MongoDB uses Perl compatible regular expressions (i.e. “PCRE”) version 8.38 with UTF-8 support.

Example: `db.bank_data.find({"first_name":{$regex:/^RI.*/}},{"first_name":1,"last_name":1,"_id":0}).pretty();`

Update and Remove Queries

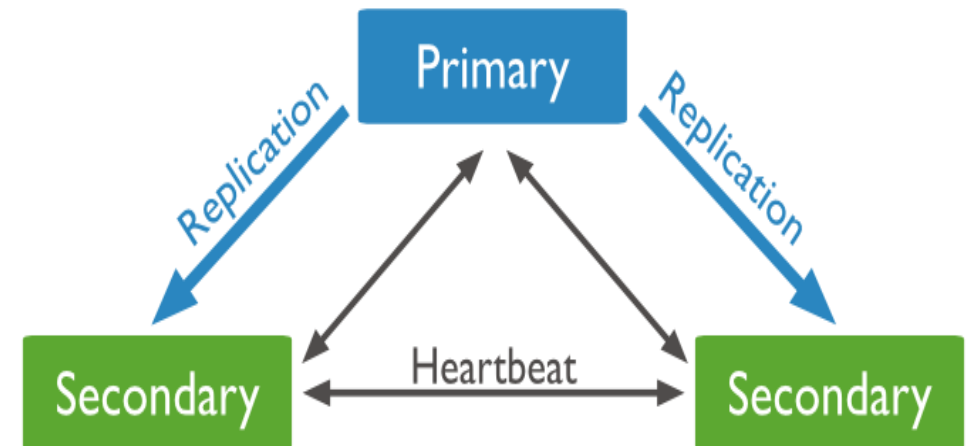
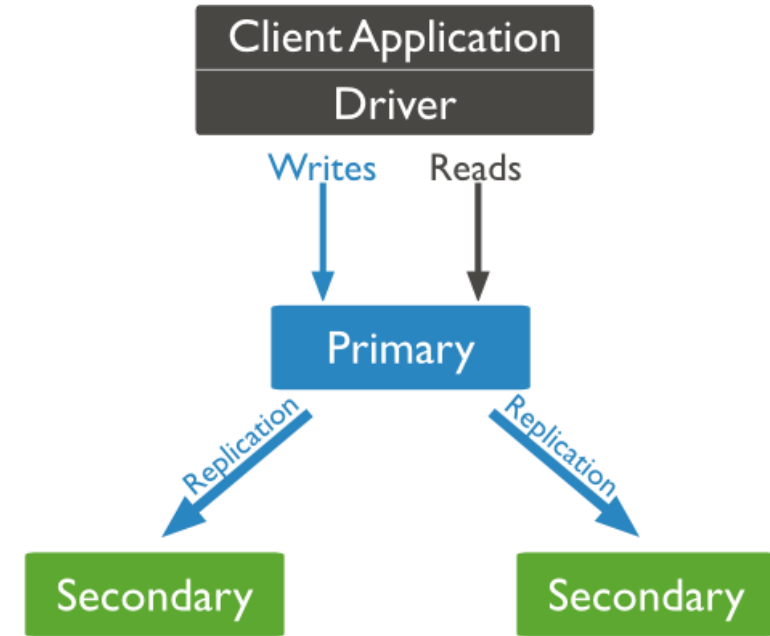
- MongoDB's **update()** and **save()** methods are used to update document into a collection. The update() method update values in the existing document while the save() method replaces the existing document with the document passed in save() method.

<code>\$rename</code>	Renames a field.
<code>\$setOnInsert</code>	Sets the value of a field if an update results in an insert of a document. Has no effect on update operations that modify existing documents.
<code>\$set</code>	Sets the value of a field in a document.
<code>\$unset</code>	Removes the specified field from a document.
<code>\$min</code>	Only updates the field if the specified value is less than the existing field value.
<code>\$max</code>	Only updates the field if the specified value is greater than the existing field value.

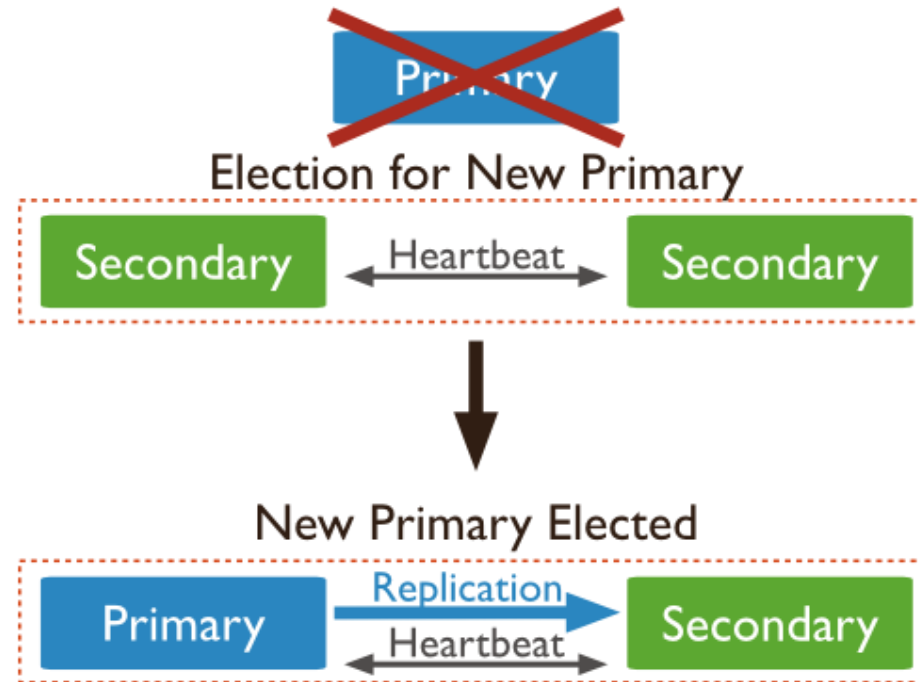
- MongoDB's **remove()** method is used to remove document from the collection. remove() method accepts two parameters. One is deletion criteria and second is justOne flag
- Example: `db.collection.remove({"last_name":"SMITH"});`

Replication Architecture


- Replication set – multiple mongod instances
- Upto 50 secondary allowed
- Primary authoritative node
 - Read
 - Writes
- Secondary nodes
 - Reads
- At most 12 replica sets are allowed.



Automatic Failover



Partition Tolerance



How does MongoDB solve the CAP?

- CP System
- Write Concern?
 - Lets you choose when was a write successful. Acknowledgements like
 - error ignored
 - How many nodes must have write acknowledged.
- Read Preferences:
 - Allows you to choose where to read from.
 - Master / Slave

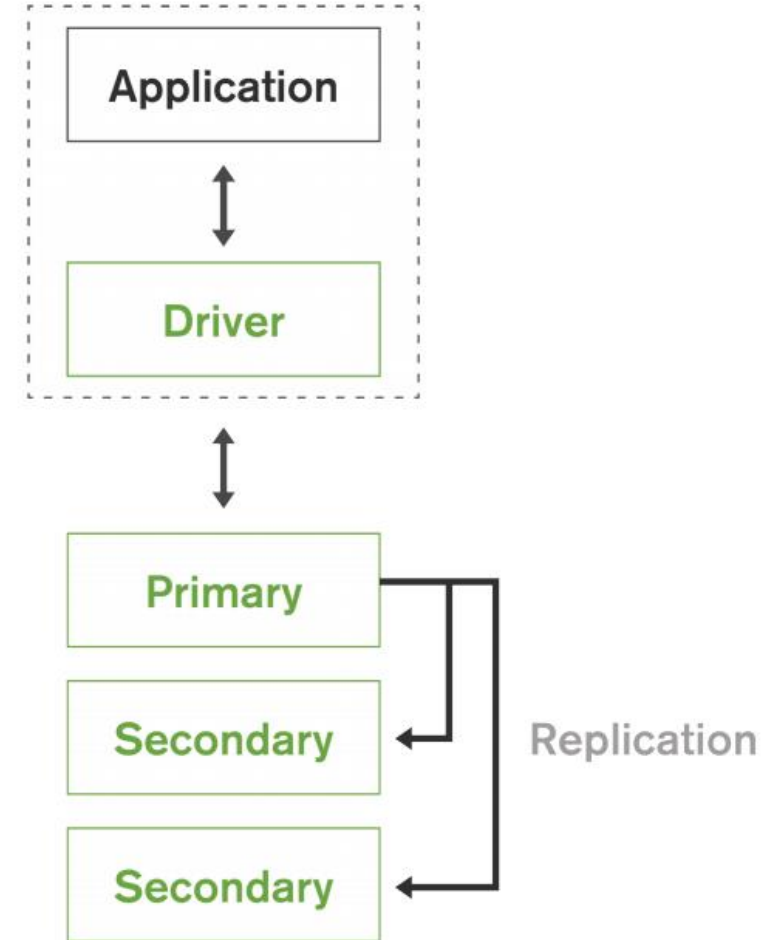


Figure 8: Self-Healing MongoDB Replica Sets for High Availability

- For Write
 - a write quorum can be implemented using **write concern**.
 - if number of available nodes < specified in write concern, write operation will fail.
 - All read operations sent to the primary are consistent to last write operations.
-

- Reads to a primary have **strict consistency**
 - Reads reflect the latest changes to the data
- Reads to a secondary have **eventual consistency**
 - Updates propagate gradually
- If clients permit reads from secondary sets – then client may read a previous state of the database.

<http://www.ccs.neu.edu/home/kathleen/classes/cs3200/20-NoSQLMongoDB.pdf>



References

- <http://jandiandme.blogspot.de/2013/06/mongodb-and-cap-theorem.html>
- <http://stackoverflow.com/questions/7339374/nosql-what-does-it-mean-for-mongodb-or-bigtable-to-not-always-be-available>
- <http://www.ccs.neu.edu/home/kathleen/classes/cs3200/20-NoSQLMongoDB.pdf>
- <https://docs.mongodb.org/manual/core/replica-set-members/>
- <https://github.com/sedouard/mongodb-mva>